

CI/CD for Microservices

Goal:

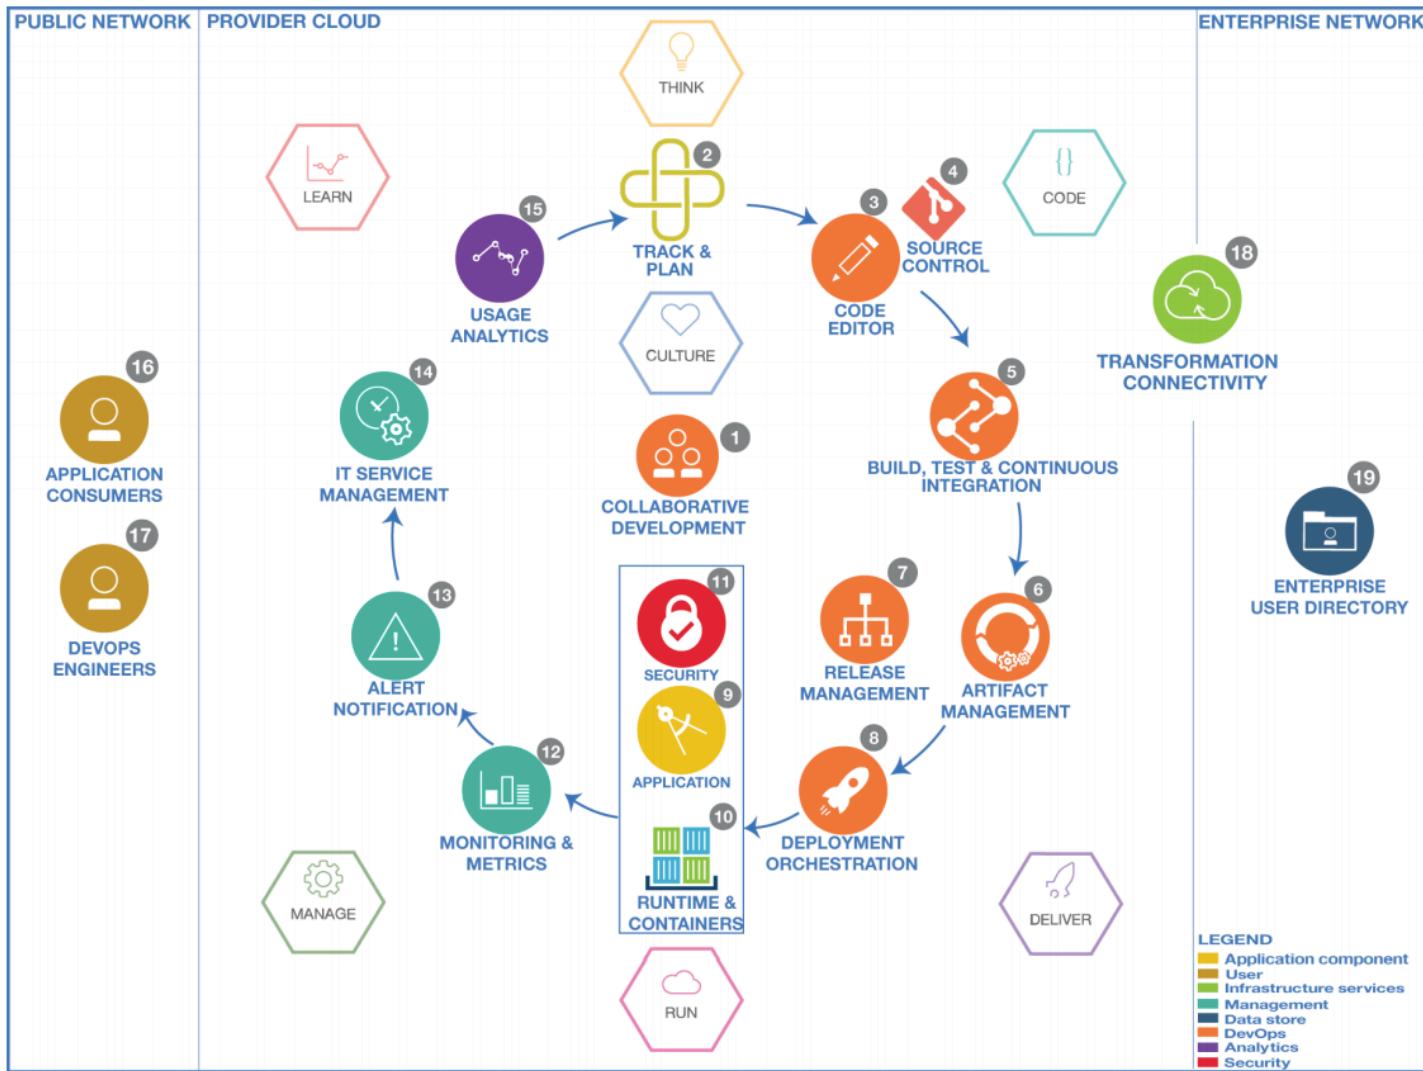
- Continuous Integration and Continuous Deployment
- For Microservices

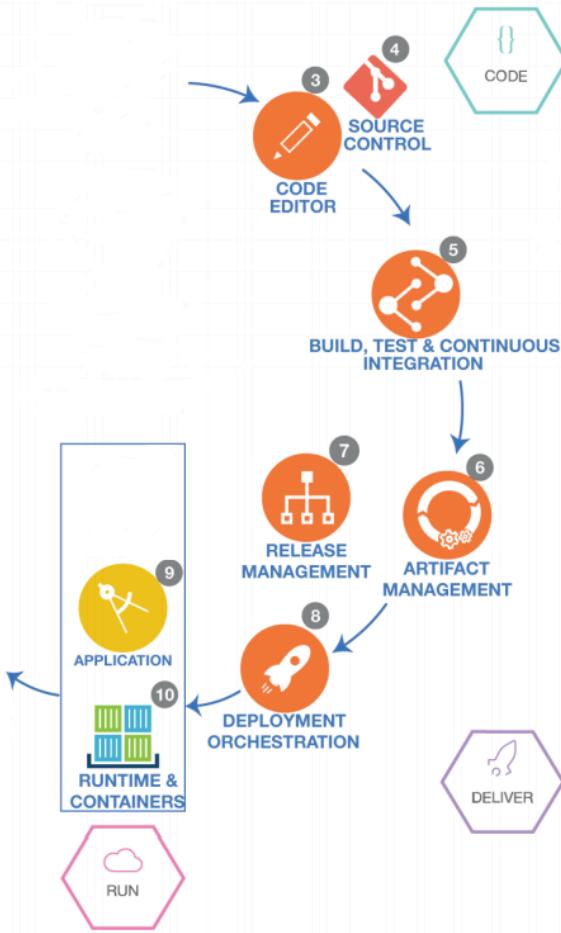
Best practices:

- Release often,
- Blue green deployment,
- Canary deployment,
- A/B Testing,
- Commit daily, reduce branching.

<https://12factor.net/>

1. Codebase: each codebase has its own CSM repo
2. Dependencies: explicitly declare and isolate dependencies
3. Config: store config in the environment, not in the application
4. Backing services: treat backing services, local or remote, as attached but swappable resources
5. Build, release, run: Strictly separate build and run stages
6. Processes: execute the app as one or more stateless processes running in the environment
7. Port binding: export services via port binding, do not rely on runtime injection
8. Concurrency: scale out via the process model, so that the application is able to span multiple processes on multiple physical machines
9. Disposability: maximize robustness with fast startup and graceful shutdown, elastic scaling, and rapid deployment
10. Dev/prod parity:
 1. Small time gap: from code to deployment asap
 2. Small personnel gap: developer is involved in dev, ops and prod
 3. Small tools gap: keep dev, staging, and prod as similar as possible
11. Logs: Treat logs as event streams, use stdout
12. Admin processes: Run admin/management tasks as one-off processes w the same rules as above





Overwhelmed? Please see the CNCF Trail Map. That and the interactive landscape are at l.cncf.io

Database

Streaming & Messaging

Application Definition & Image Build

Continuous Integration & Delivery

Platform

Observability and Analysis

App Definition and Development



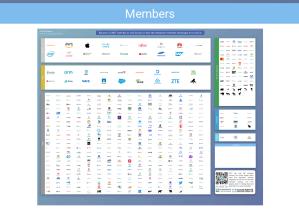
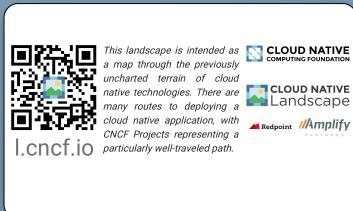
Orchestration & Management



Runtime

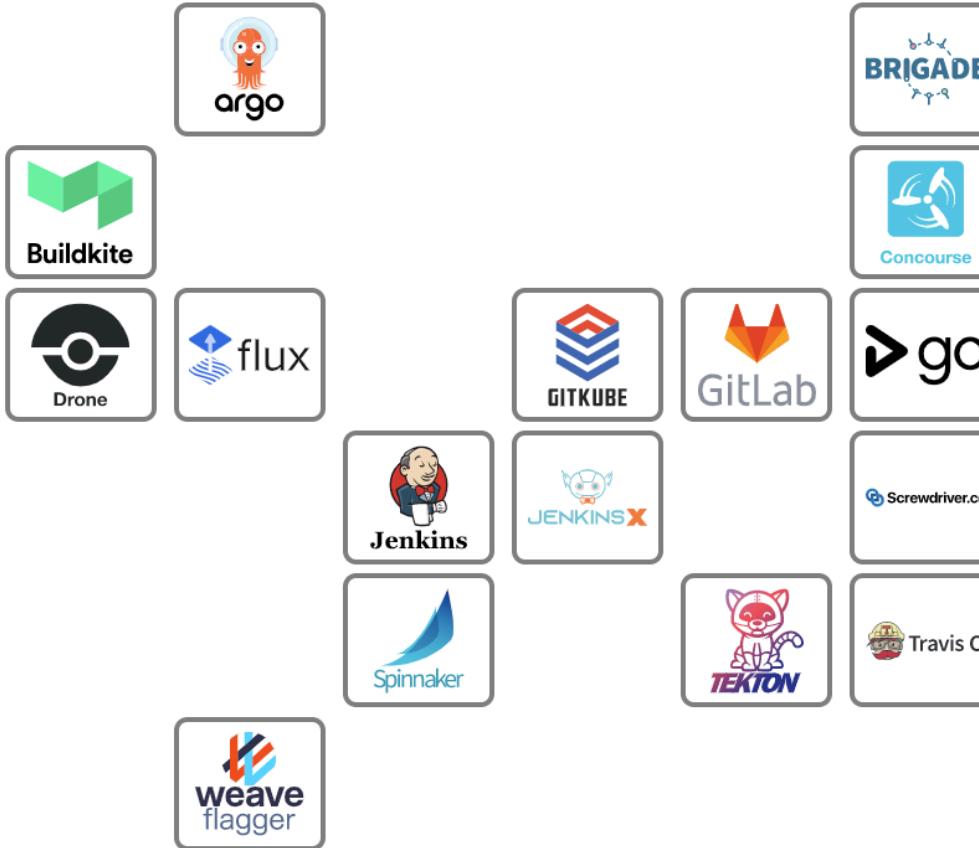


Provisioning



Continuous Integration & Delivery

<https://landscape.cncf.io/license=open-source>



Application Definition & Image Build

Scheduling & Orchestration

Service Mesh



CNCF Incubating



Buildpacks.io



DevSpace



Octant



okteto



kubernetes

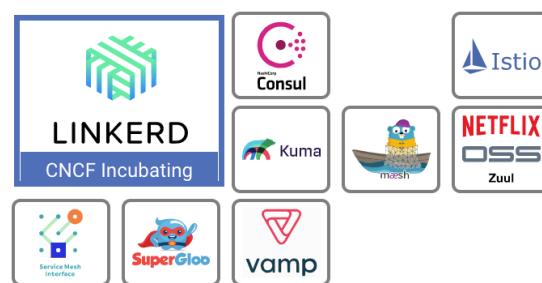
CNCF Graduated



MESOS



Crossplane



LINKERD

CNCF Incubating



CIT



Build, Test, Run

Apache Maven

```
$ brew install maven
$ mvn -B archetype:generate -DarchetypeGroupId=org.apache.maven.archetypes
-DgroupId=com.remkohdev.app -DartifactId=my-app
$ cd my-app
$ mvn clean install
$ mvn test
$ java -cp target/my-app-1.0-SNAPSHOT.jar com.remkohdev.app.App
```

Spring Framework

```
$ brew tap pivotal/tap
$ brew install springboot
$ spring init --dependencies=web,data-rest,thymeleaf MyApp
$ cd MyApp
$ mvn clean install
$ mvn test
$ mvn spring-boot:run
```

Node.js - Loopback

```
$ npm install -g loopback-cli
$ lb
? What's the name of your application? MyNodeApp
? Which version of LoopBack would you like to use? 3.x (current)
? What kind of application do you have in mind? api-server
$ cd MyNodeApp
$ npm install
$ npm test
$ npm start
```

Test

Spring JUnit

```
@RunWith(SpringRunner.class)
@WebMvcTest(value = APIController.class, secure = false)
public class APIControllerTest {
    @Autowired
    private MockMvc mockMvc;
    @Test
    public void getMessage() throws Exception {

        String name = "venus";
        RequestBuilder requestBuilder = MockMvcRequestBuilders
            .get("/api/hello?name="+name)
            .accept(MediaType.APPLICATION_JSON);
        MvcResult result = (MvcResult) mockMvc.perform(requestBuilder)
            .andExpect(status().isOk())
            .andExpect(jsonPath("$.message", is("Hello "+name)))
            .andDo(print())
            .andReturn();
        System.out.println(result.getResponse());
    }
}
```

Test

Mochajs

```
let request = require( 'supertest' )
let assert = require( 'chai' ).assert
let serverPromise = require( '../server/server' )

describe( 'CRUD', () => {
  let server = null
  before( async function () {
    server = await serverPromise
  })
  it( 'GET /api/hello', function ( done ) {
    request( server ).get( '/v1.1/banks' )
      .set( 'Content-type', 'application/json' )
      .expect( 200 )
      .then( response => {
        assert.isOk( response.body.length > 0 )
        done()
      })
  })
})
```

Test

Postman

```
$ curl -X GET 'http://localhost:8080/api/hello?name=pluto'
```

```
pass = true;
fail = false;
try {
    resp = JSON.parse(pm.response.text());
    pm.test("Check response to have status 200", function () {
        pm.response.to.have.status(200);
    });
    pm.test("Check response includes message", function () {
        pm.expect(pm.response.text()).to.include("message");
    });
    pm.globals.set("env_message", resp["message"]);
}
catch(err) {
    tests["Parse response JSON or tests initialisation"] = fail;
}
```

```
$ newman run postman/springclient-test.postman_collection.json
```

Run

Apache Maven

```
$ mvn clean install  
$ java -cp target/my-app-1.0-SNAPSHOT.jar com.remkohdev.app.App
```

Spring Framework

```
$ mvn clean install  
$ mvn spring-boot:run
```

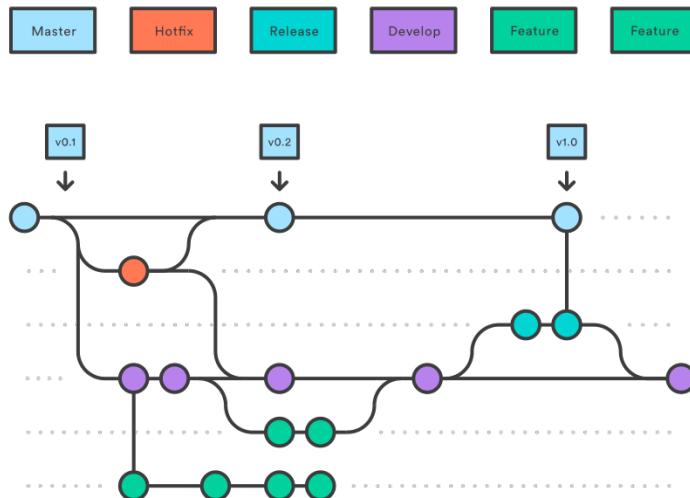
Node.js - Loopback

```
$ npm install  
$ npm start
```

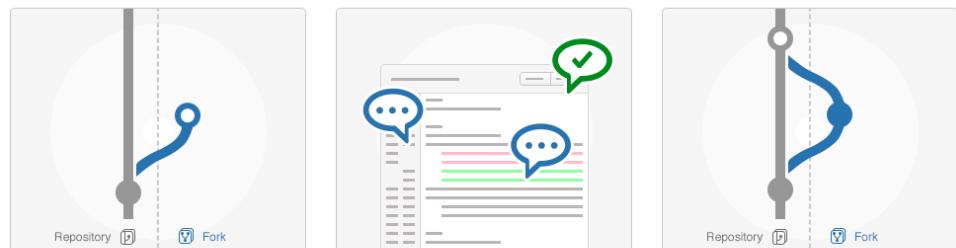
Git

```
$ git add .
$ git commit -m 'my update for my issue'
$ git push
```

Feature branches



Forks



Fork

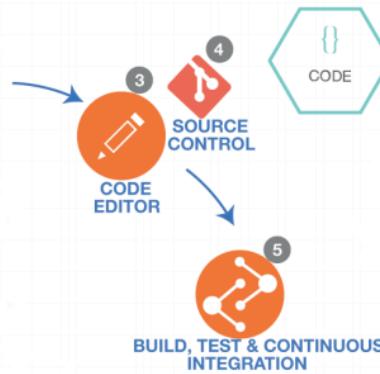
Develop features on a branch and create a pull request to get changes reviewed.

Discuss

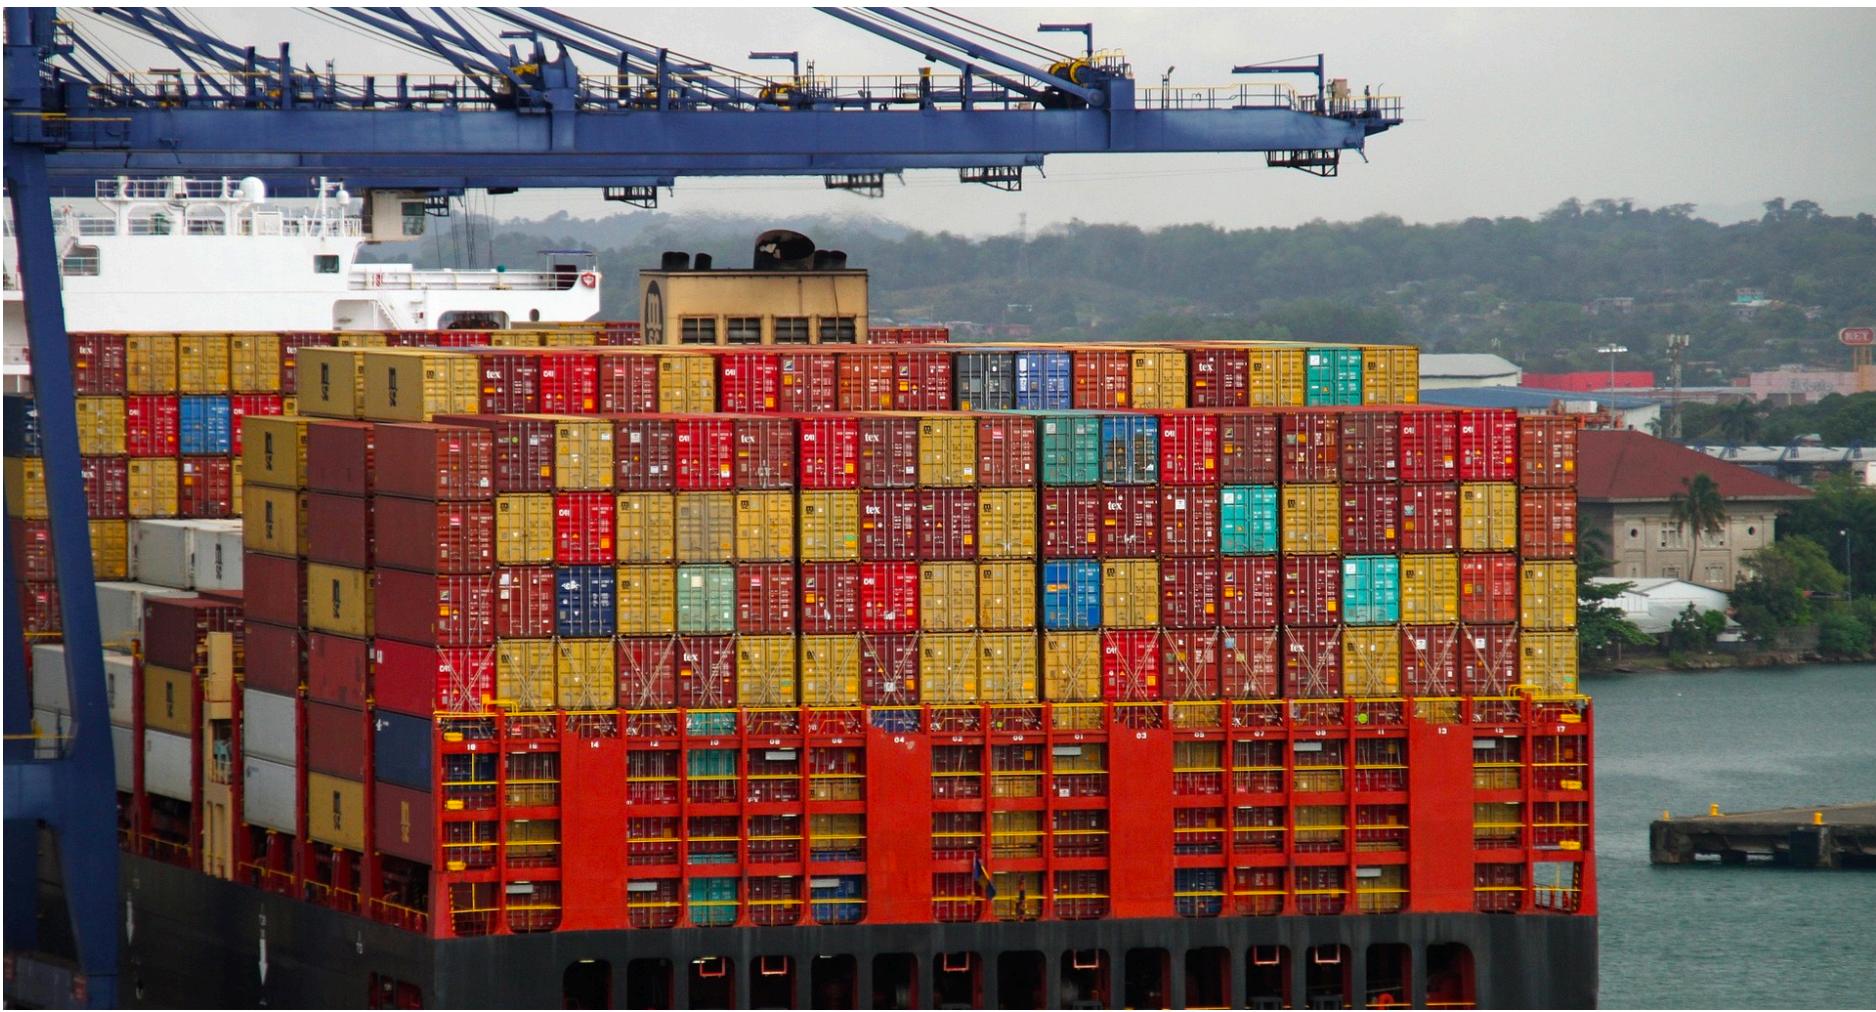
Discuss and approve code changes related to the pull request.

Merge

Merge the branch with the click of a button.



Docker



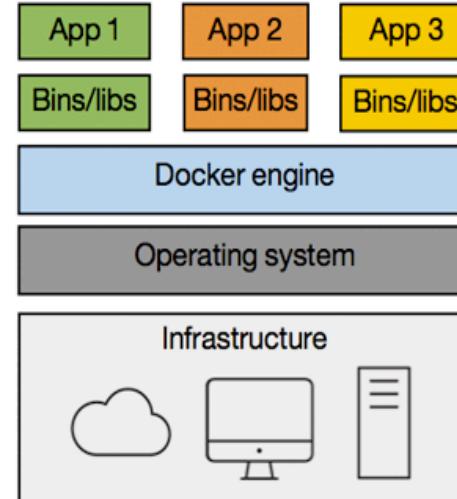
Dockerfile

```
FROM python:3.6.1-alpine  
RUN pip install flask  
CMD ["python","app2.py"]  
COPY app2.py /app2.py
```

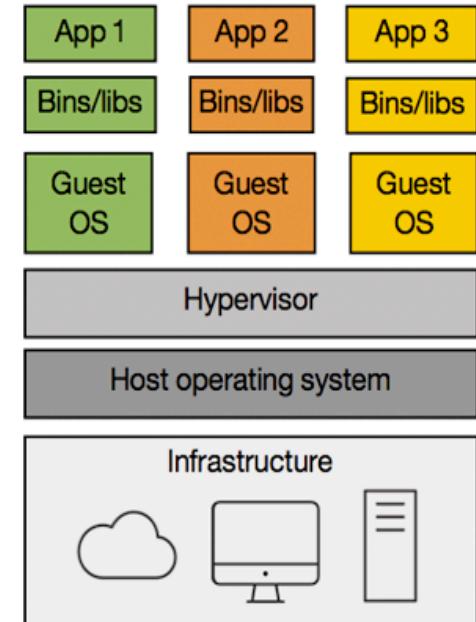
Image Layers



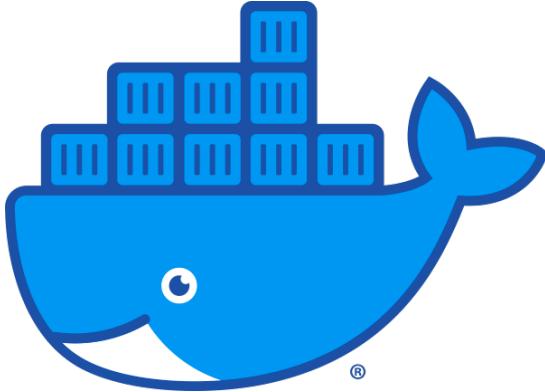
Deployment through Containers



Containers



Virtual Machines



Dockerfile

```
FROM node:11.10.1-stretch-slim

ARG NPM_SCOPE
ARG NPM_REGISTRY_URL
ARG ART_PASSWORD
ARG ART_USERNAME

# Create app directory
WORKDIR /usr/src/app

# Bundle app source
COPY ./package*.json .
COPY ./npm-shrinkwrap.json .
# access private npm registry
RUN echo '@'$NPM_SCOPE':registry=https://'$NPM_REGISTRY_URL'\n//'$NPM_REGISTRY_URL':_password=''$ART_PASSWORD'\n//'$NPM_REGISTRY_URL':username=''$ART_USERNAME'\n//'$NPM_REGISTRY_URL':email=''$ART_USERNAME'\n//'$NPM_REGISTRY_URL':always-auth=true'>.npmrc

RUN npm install
RUN npm audit fix

RUN rm -f .npmrc

COPY ..

EXPOSE 3000

CMD [ "npm", "start" ]
```

docker-compose.yml

```
version: '3'
services:
  my-api:
    build:
      context: .
      dockerfile: ./Dockerfile
    args:
      NPM_SCOPE: '${NPM_SCOPE}'
      NPM_REGISTRY_URL: '${ART_NPM_REGISTRY_URL}'
      ART_PASSWORD: '${ART_BASE64_PASSWORD}'
      ART_USERNAME: '${ART_USERNAME}'
    command: npm start
    ports:
      - "3000:3000"
    environment:
      NODE_ENV: local
      COUCHDB_USERNAME: ${COUCHDB_USERNAME}
      COUCHDB_PASSWORD: ${COUCHDB_PASSWORD}
      COUCHDB_SERVER: ${COUCHDB_SERVER}
      COUCHDB_PORT: ${COUCHDB_PORT}
      COUCHDB_PROTOCOL: ${COUCHDB_PROTOCOL}
    volumes:
      - .:/usr/src/app

  networks:
    - my-network
```

Docker

```
$ docker build --no-cache -t my-app .
$ docker tag my-app:latest remkohdev/my-app:0.1.0
$ docker run -d --name my-app -p 3000:3000 my-app
$ docker login
$ docker push remkohdev/my-app:0.1.0
```

Docker Compose

```
$ docker-compose build
$ docker-compose up -d
```

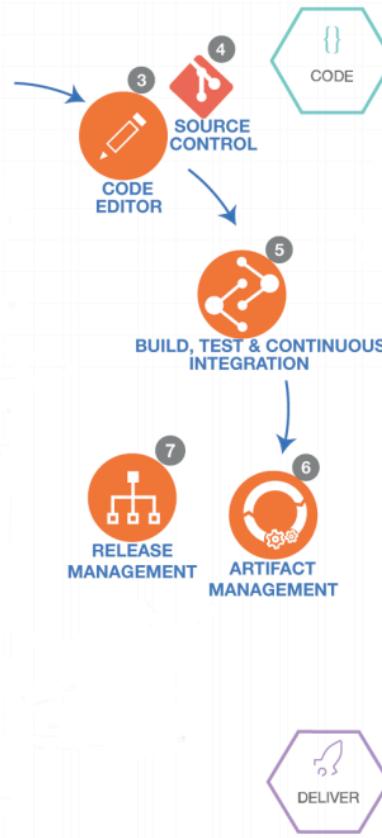


CLOUD FOUNDRY

manifest.yml

```
---
applications:
- name: my-app
  buildpack: https://github.com/cloudfoundry/nodejs-buildpack
  command: node my-app.js
```

```
$ cf push my-app -c "node my-app.js"
```





We can build, test and run on localhost

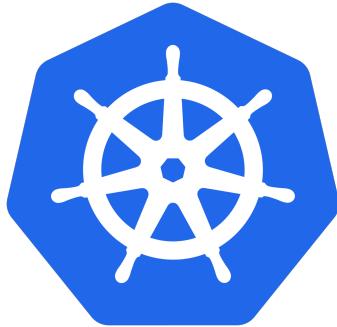
We can collaborate

We can run anywhere

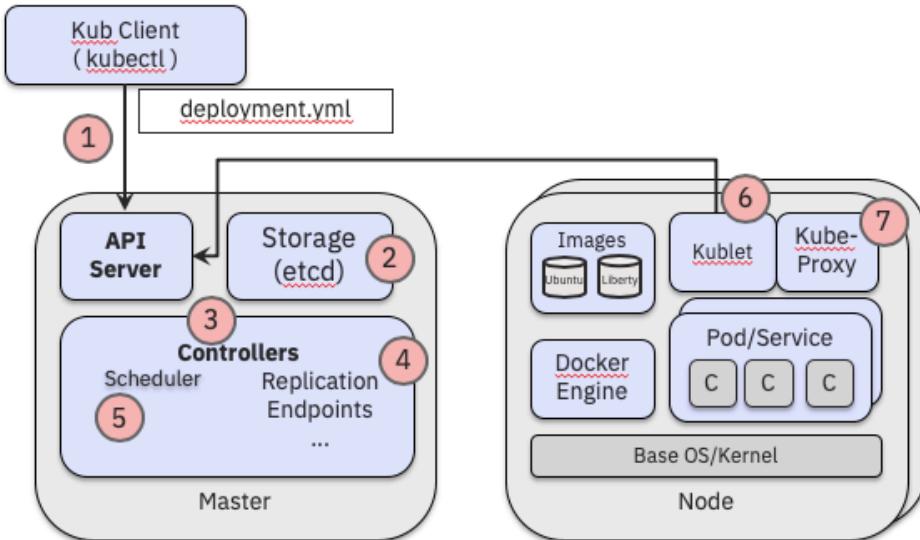
We pushed our build artifact

Let's deploy

Kubernetes



Kubernetes is a container-orchestration system for automating application deployment, scaling, and management, originally created by Google, now maintained by the Cloud Native Computing Foundation (CNCF).



1. User via "kubectl" deploys a new application
2. API server receives the request and stores it in the DB (etcd)
3. Watchers/controllers detect the resource changes and act upon it
4. ReplicaSet watcher/controller detects the new app and creates new pods to match the desired # of instances
5. Scheduler assigns new pods to a kubelet
6. Kubelet detects pods and deploys them via the container running (e.g. Docker)
7. KubeProxy manages network traffic for the pods – including service discovery and load-balancing

Kubernetes Objects

<https://kubernetes.io/docs/reference/generated/kubernetes-api/v1.16/>

Basic Objects:

- Pod
- Service
- PersistentVolume
- PersistentVolumeClaim
- Namespace
- ConfigMap
- Secret

Higher-Level Abstractions:

- Deployment
- DaemonSet
- StatefulSet
- ReplicaSet
- Job

Custom Resource Definitions (CRD):

- Operators (<https://operatorhub.io/>)

deployment.yaml

```
apiVersion: apps/v1beta2
kind: Deployment
metadata:
  name: my-app-deployment
  namespace: my-ns
  labels:
    app: my-app
spec:
  replicas: 1
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: my-app
          image: remkohdev/my-app:latest
          ports:
            - name: main
              protocol: TCP
              containerPort: 6661
          envFrom:
            - configMapRef:
                name: my-app-configmap
  resources:
    requests:
      memory: "120M"
      cpu: "500m"
```

Deploy

```
# configmap
kubectl delete configmap -n my-ns my-app-configmap

kubectl create -f ./helm/templates/configmap.yaml

# deployment
kubectl delete deployment -n my-ns my-app-deployment

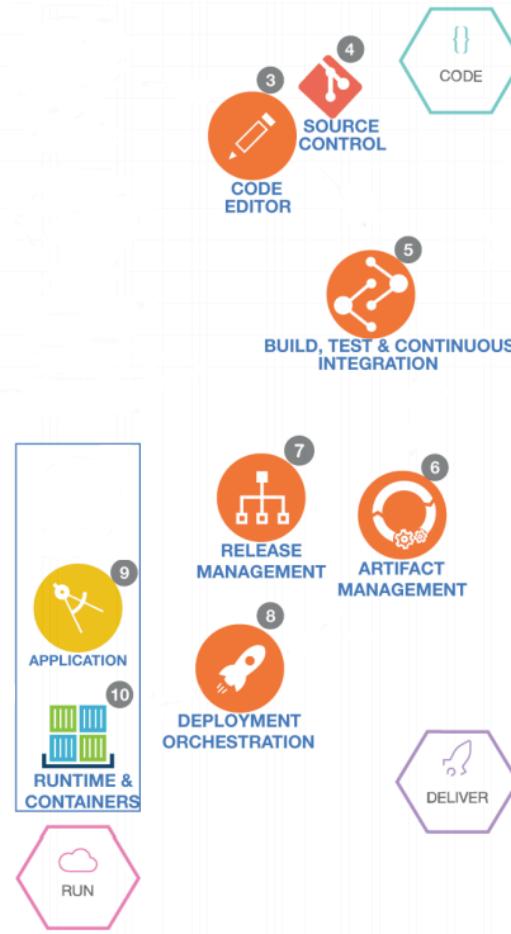
# while resource still exists wait
rc=$(eval 'kubectl get deployment -n my-ns my-app-deployment')
while [ ! -z "$rc" ]
do
    rc=$(eval 'kubectl get deployment -n my-ns my-app-deployment')
done

kubectl create -f ./helm/templates/deployment.yaml

# service
kubectl delete svc -n my-ns my-app-svc

# while resource still exists wait
rc=$(eval 'kubectl get svc -n my-ns my-app-svc')
while [ ! -z "$rc" ]
do
    rc=$(eval 'kubectl get svc -n my-ns my-app-svc')
done

kubectl create -f ./helm/templates/svc.yaml
```



Automation

- Bash
- Docker-compose
- Jenkins

```
docker build --no-cache -t my-app .
docker tag my-app:latest remkohdev/my-app:0.1.0
docker login -u remkohdev
docker push remkohdev/my-app:0.1.0
```

```
kubectl delete configmap -n my-ns my-app-configmap
kubectl create -f ./helm/templates/configmap.yaml

kubectl delete deployment -n my-ns my-app-deployment
rc=$(eval 'kubectl get deployment -n my-ns my-app-deployment')
while [ ! -z "$rc" ]
do
    rc=$(eval 'kubectl get deployment -n my-ns my-app-deployment')
done
kubectl create -f ./helm/templates/deployment.yaml

kubectl delete svc -n my-ns my-app-svc
rc=$(eval 'kubectl get svc -n my-ns my-app-svc')
while [ ! -z "$rc" ]
do
    rc=$(eval 'kubectl get svc -n my-ns my-app-svc')
done
kubectl create -f ./helm/templates/svc.yaml
```

Jenkins

```
node() {  
  
    stage("Checkout SCM") {  
        checkout scm;  
    }  
  
    stage("Artifactory Configuration") {  
        server = Artifactory.server( 'na.artifactory.swg-devops.com' )  
        buildInfo = Artifactory.newBuildInfo()  
        rtDocker = Artifactory.docker server: server  
        env.ART_DOCKER_IMAGE_TAG = "${ART_DOCKER_REPO}/${ART_B  
UILD_BRANCH}/${DOCKER_IMAGE_NAME}:${env.ART_DOCKER_TAG}"  
    }  
  
    stage("Node Test"){  
        sh 'node -v'  
    }  
  
    stage("Build Docker Image") {  
        docker.build( "${ART_DOCKER_IMAGE_TAG}" )  
    }  
  
    stage("Push Docker Image") {  
        withCredentials([usernamePassword(credentialsId: "$JENKINS_ART  
IFACTORY_CREDENTIAL", passwordVariable: 'ART_PW', usernameVaria  
ble: 'ART_USER')]) {  
            rtDocker.push( "${ART_DOCKER_IMAGE_TAG}", "${ART_DOCKER_  
REPO}" )  
        }  
    }  
}
```

Jenkins

```
stage("Set KUBE.Context") {
    withCredentials([usernamePassword(credentialsId: "$JENKINS_IKS_CREDENTIALS", passwordVariable: 'IKS_PW', usernameVariable: 'IKS_USER')]) {
        sh """
            mkdir ./tmp
            curl -sL https://ibm.biz/idt-installer | bash
            ibmcloud login --skip-ssl-validation -u ${IKS_USER} -p ${IKS_PW} -a https://${KUBE_URL} -c my-account
        """
    }
}

stage("Deploy to KUBE") {
    sh """
        chmod 777 ${DEPLOYMENT_FILE}
        ./${DEPLOYMENT_FILE}
    """
}
```

Github Webhooks

remkohdev / spring-client

Unwatch 1 Star 0 Fork 1

Code Issues 2 Pull requests 1 ZenHub Projects 0 Wiki Security Insights Settings

Webhooks

Add webhook

Webhooks allow external services to be notified when certain events happen. When the specified events happen, we'll send a POST request to each of the URLs you provide. Learn more in our [Webhooks Guide](#).

✓ <https://jenkins-ci.cda Openshift Cluster-1c0e8fb1c68214cf875a9ca7dd1e060-0001.us-south.co...>
(pull_request and push)

Edit Delete

Options
Collaborators
Branches
Webhooks
Notifications
Integrations & services
Deploy keys
Autolink references

Jenkins Pipeline

✓ spring-client < 4

Pull Request: PR-2 [PR-2](#) 1m 44s Commit: 0dd5fa0 3 days ago No changes Restarted from build #3, stage Login

Pipeline Changes Tests Artifacts [Logout](#)



Create Project - 4s

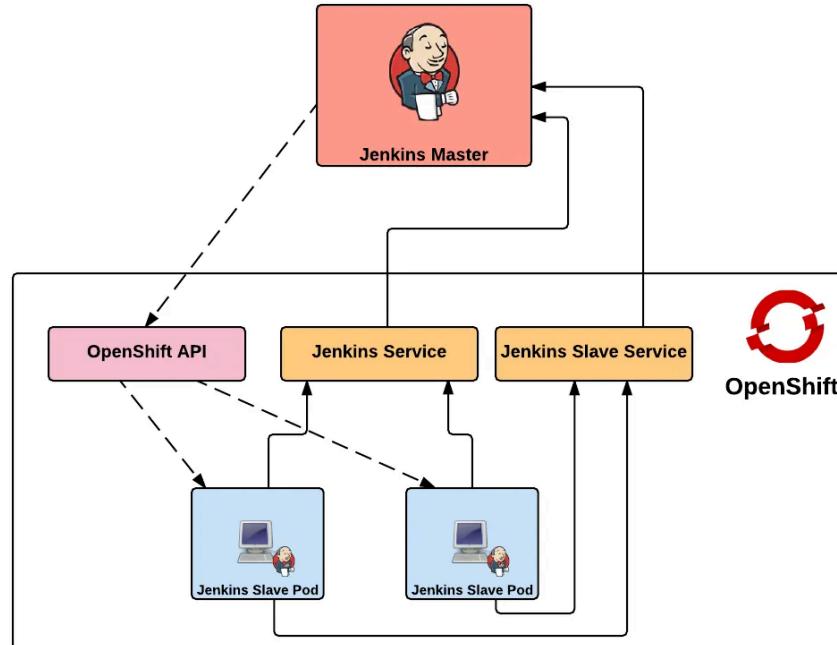
✓	> Create Project — Print Message	<1s
✓	> oc new-project springclient-ns — Shell Script	1s
✓	> oc project springclient-ns — Shell Script	<1s
✓	> Print Message	2s
✓	> Using project: null — Print Message	<1s

Jenkins external to OpenShift

When leveraging the Kubernetes plugin for Jenkins, Jenkins runs dynamic agents in a Kubernetes cluster to dynamically provision slave instances. There must be communication between the Jenkins master and OpenShift:

- Jenkins master communicates with the OpenShift API to manage the lifecycle of slave instances and take advantage of the elasticity OpenShift provides.
- OpenShift and Jenkins communicate via the port exposed by the web console
- Jenkins master communicates via a TCP port used for the JNLP slaves

<https://github.com/jenkinsci/kubernetes-plugin>







We can build, test and run on localhost

We can collaborate

We can run anywhere

We pushed our build artifact

We deployed

We automated

We have a CI/CD for Microservices

Okay I knew that...

Plus,

That's a lot more workload than I used to have, before Microservices

Kubernetes Extensions

Extending Kubernetes

Customization can be divided into:

- configuration,
- Extensions.

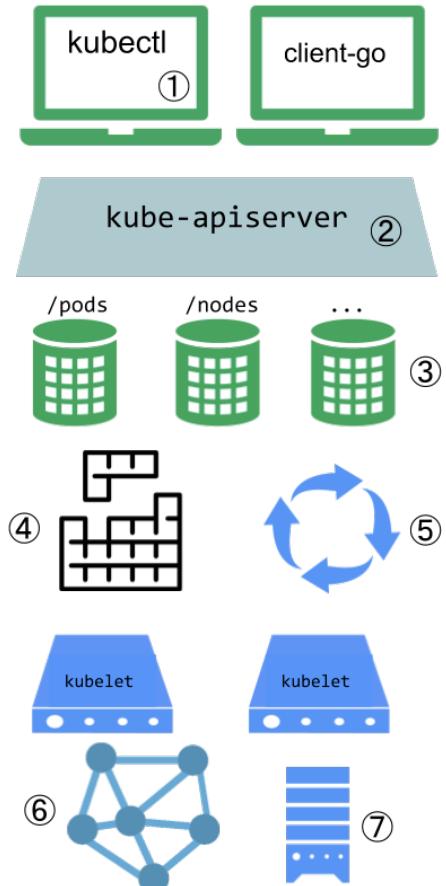
You can extend the Kubernetes by using the Controller pattern or the webhook model.

Controllers read a “spec”, do work, and then update the “status”.

When Kubernetes is the client that calls out to a remote service, it is called a Webhook and the remote service is the Webhook Backend.

Extension points:

1. Kubectl, kubectl plugins,
2. Extensions in the apiserver allow authenticating or blocking requests, editing content and handling deletion,
3. Custom Resources (CR) using CustomResourceDefinition API and Operators,
4. Scheduler Extensions,
5. Controllers are often used with CR,
6. Node-level Network Plugins, CNI Plugins or Kubenet plugins,
7. Storage Plugins,



Custom Resource Definition (CRD)

In June 2017, Kubernetes v1.7 introduced Custom Resource Definitions (CRD).

A *Custom Resource (CR)* is an extension of the Kubernetes API that allows you to store an API object of a kind. Even many core Kubernetes functions are now built using custom resources.

A *Custom Resource Definition (CRD)* file defines your own object kinds and lets the API Server handle its lifecycle.

Example: CRD

guestbook.yaml

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: guestbooks.apps.ibm.com
spec:
  group: apps.ibm.com
  versions:
    - name: v1
      served: true
      storage: true
      schema:
        openAPI3Schema:
          type: object
          properties:
            spec:
              type: object
              properties:
                guestbookTitle:
                  type: string
                guestbookSubtitle:
                  type: string
  scope: Namespaced
  names:
    plural: guestbooks
    singular: guestbook
    kind: Guestbook
    shortNames:
      - gb
```

my-guestbook.yaml

```
apiVersion: "apps.ibm.com/v1"
kind: Guestbook
metadata:
  name: my-guestbook
spec:
  guestbookTitle: "The Chemical Wedding of Remko"
  guestbookSubtitle: "First Day"
```

```
% kubectl create -f guestbook-crd.yaml
customresourcedefinition.apiextensions.k8s.io/guestbooks.apps.ibm.com created
% kubectl create -f my-guestbook.yaml
guestbook.apps.ibm.com/my-guestbook created
% kubectl get guestbooks
NAME      AGE
my-guestbook  3m51s
% kubectl describe guestbook my-guestbook
Name:      my-guestbook
Namespace: default
Labels:    <none>
Annotations: <none>
API Version: apps.ibm.com/v1
Kind:      Guestbook
Metadata:
  Creation Timestamp: 2020-04-20T05:40:46Z
  Generation: 1
  Resource Version: 295507
  Self Link:   /apis/apps.ibm.com/v1/namespaces/default/guestbooks/my-guestbook
  UID:        e683caef-33f5-4f83-b54e-85ee163e0abf
Spec:
  Guestbook Subtitle: First Day
  Guestbook Title: The Chemical Wedding of Remko
Events:    <none>
```



≡ Custom Resource Definitions > guestbooks.apps.ibm.com

Storage Classes

Namespace

default

Overview

Workloads

Cron Jobs

Daemon Sets

Deployments

Jobs

Pods

Replica Sets

Replication Controllers

Stateful Sets

Discovery and Load Balancing

Ingresses

Services

Config and Storage

Config Maps

Persistent Volume Claims

Secrets

Custom Resource Definitions

Metadata

Name

guestbooks.apps.ibm.com

Creation time

Apr 20, 2020

Age

17 minutes

UID

2dc92a69-1e5e-489f-a9fa-c5e7af780aab

Resource Information

Version

v1

Scope

Namespaced

Group

apps.ibm.com

Accepted Names

Plural

guestbooks

Singular

guestbook

Kind

Guestbook

List Kind

GuestbookList

Short Names

gb

Objects

Name

Namespace

Age

my-guestbook

default

10 minutes

1 – 1 of 1

Versions

Name

Served

Storage

v1

True

True



Operator Pattern

Custom Resources alone let you store and retrieve structured data. The [Operator pattern](#) combines custom resources and custom controllers. Combined with a *custom controller*, custom resources provide a true *Declarative API*. You can use custom controllers to encode domain knowledge for specific applications into an extension of the Kubernetes API.

Consider API aggregation if:

Your API is [Declarative](#).

Read and write your new types using kubectl.

View your new types in a Kubernetes UI, such as dashboard, alongside built-in types.

You are developing a new API.

Accept the format restriction that Kubernetes puts on REST resource paths, such as API Groups and Namespaces. (See the [API Overview](#).)

Your resources are naturally scoped to a cluster or namespaces of a cluster.

You want to reuse [Kubernetes API support features](#).

Prefer a stand-alone API if:

Your API does not fit the [Declarative](#) model.

kubectl support is not required

Kubernetes UI support is not required.

You already have a program that serves your API and works well.

You need to have specific REST paths to be compatible with an already defined REST API.

Cluster or namespace scoped resources are a poor fit; you need control over the specifics of resource paths.

You don't need those features.

Operator Framework SDK

```
operator-sdk new $OPERATOR_PROJECT --type go --repo github.com/$DOCKER_USERNAME/$OPERATOR_NAME
operator-sdk add api --api-version=$OPERATOR_GROUP/$OPERATOR_VERSION --kind=$CRD_KIND
operator-sdk add controller --api-version=$OPERATOR_GROUP/$OPERATOR_VERSION --kind=$CRD_KIND
operator-sdk build docker.io/$DOCKER_USERNAME/$OPERATOR_NAME
docker login docker.io -u $DOCKER_USERNAME
docker push docker.io/$DOCKER_USERNAME/$OPERATOR_NAME
sed -i "s|REPLACE_IMAGE|docker.io/$DOCKER_USERNAME/$OPERATOR_NAME|g" deploy/operator.yaml
oc create sa $OPERATOR_PROJECT
oc create -f deploy/role.yaml
oc create -f deploy/role_binding.yaml
oc create -f deploy/crds/${OPERATOR_GROUP}_${CRD_KIND,,}s_crd.yaml
oc create -f deploy/operator.yaml
oc create -f deploy/crds/${OPERATOR_GROUP}_${OPERATOR_VERSION}_${CRD_KIND,,}_cr.yaml
oc get deployment $OPERATOR_PROJECT
oc get pod -l app=example-${CRD_KIND,,}
oc describe ${CRD_KIND,,}s.${OPERATOR_GROUP} example-${CRD_KIND,,}
```

Operator Framework SDK

```
// Reconcile reads state of the cluster for a Guestbook object and makes changes based on the state  
// read and what is in the Guestbook.Spec  
// TODO(user): User must modify this Reconcile function to implement their own Controller logic.  
This example creates a Pod as an example  
func (r *ReconcileGuestbook) Reconcile(request reconcile.Request) (reconcile.Result, error) {  
    ...  
    // Fetch the Guestbook instance  
    instance := &guestbookv1.Guestbook{  
        ...  
        // Define a new Pod object  
        pod := newPodForCR(instance)  
        ...  
    }  
}
```

```
guestbook-project  
> build  
> cmd  
< deploy  
< crds  
  ! guestbook.remkoh.dev_guestbooks_crd.yaml  
  ! guestbook.remkoh.dev_v1_guestbook_cr.yaml  
  ! operator.yaml  
  ! role.yaml  
  ! role_binding.yaml  
  ! service_account.yaml  
< pkg  
< apis  
< guestbook  
< v1  
  -> doc.go  
  -> guestbook_types.go  
  -> register.go  
  -> zz_generated.deepcopy.go  
  -> group.go  
  -> addtoscheme_guestbook_v1.go  
  -> apis.go  
< controller  
  -> controller.go  
> version  
❖ .gitignore  
≡ go.mod
```

Application (CRD)

<https://github.com/kubernetes-sigs/application>

```
apiVersion: app.k8s.io/v1beta1
kind: Application
metadata:
  name: "guestbook"
  labels:
    app.kubernetes.io/name: "guestbook"
spec:
  selector:
  matchLabels:
    app.kubernetes.io/name: "guestbook"
  componentKinds:
    - group: v1
      kind: Deployment
    - group: v1
      kind: Service
  descriptor:
    type: "guestbook"
  keywords:
    - "gb"
    - "guestbook"
  links:
    - description: Github
      url: "https://github.com/IBM/guestbook"
  version: "0.1.0"
  description: "The Guestbook application is an example app to demonstrate key
  Kubernetes functionality."
  maintainers:
    - name: IBM Developer
      email: developer@ibm.com
  owners:
    - name: IBM Developer
      email: developer@ibm.com
```

Build from Source

Source-to-Image (S2I)

Source-to-Image (S2I) builds reproducible container images from source code:

- From a builder image,
- Inject source code into a known directory,
- Create a runnable setup from source code,
- Commit the new container and set the image entrypoint.

For compiled languages, S2I creates a builder image and a runtime image.

Deploying a Spring Boot Application to Kubernetes

```
$ docker pull registry.access.redhat.com/redhat-openjdk-18/openjdk18-openshift
$ s2i build . registry.access.redhat.com/redhat-openjdk-18/openjdk18-openshift s2i-spring-client
$ oc project s2i-spring-client-ns
$ oc new-app --name springclient registry.access.redhat.com/redhat-openjdk-18/openjdk18-openshift~
  --strategy=source --allow-missing-images --build-env='JAVA_APP_JAR=hello.jar'
$ oc expose svc/springclient
$ curl -X GET 'http://<openshift-url>/api/hello?name=pluto'
```

Add Webhook

Get Github webhook

```
$ oc status
In project s2i-spring-client-ns on server https://c100-e.us-south.containers.cloud.ibm.com:30403
http://springclient-s2i-spring-client-ns.cda Openshift-cluster-9c1e3fb0c12345cf678a9ca7dd1e060-0001.us-south.
containers.appdomain.cloud to pod port 8080-tcp (svc/springclient)
dc/springclient deploys istag/springclient:latest <-
bc/springclient source builds https://github.com/remkohdev/spring-client.git#master on istag/openjdk18-openshift:latest
deployment #1 deployed 24 hours ago - 1 pod

$ oc describe bc/springclient | grep -i webhook
...
Webhook GitHub:
  URL: https://c100-e.us-south.containers.cloud.ibm.com:30403/apis/build.openshift.io/v1/namespaces/s2i-spring-client-ns/
buildconfigs/springclient/webhooks/<secret>/github
Webhook Generic:
```

You can find the URL secret for the Github webhook as a part of the full Route path,
e.g. `9c1e3fb0c12345cf678a9ca7dd1e060-0001`

Open Toolchain (OTC) Templates

<https://github.com/open-toolchain/sdk/wiki>

 Develop a Cloud Foundry app
Continuously deliver a Cloud Foundry app with repos and issue tracking hosted by IBM.

 Develop a Kubernetes app
Continuously deliver a secure Docker app to a Kubernetes Cluster.

 Develop a Kubernetes app with Razee
Continuously deliver a secure Docker app to a Kubernetes Cluster using Razee

 Develop a Kubernetes app with Helm
Continuously deliver a secure Docker app to a Kubernetes Cluster using a Helm Chart.

Build, Test, and Deploy Templates

 Develop and test a Cloud Foundry app
Continuously deliver a Cloud Foundry app.

 Develop and test microservices on Cloud Foundry
Continuously deliver a microservices app with repos and issue tracking.

 Canary testing in Kubernetes using Istio
Develop and canary test new features in your Kubernetes application using Istio.

 Garage Method tutorial with Cloud Foundry
Apply practices and tools across the DevOps lifecycle.

[Overview](#)[Connections](#)[Manage](#)

Toolchains /



toolchain-kube-guestbook

Resource Group: Default

Location: Washington DC

[Add tags](#)[Add a Tool](#)

THINK

Issues
guestbook

Configured

CODE

GitHub
guestbook

Configured

DELIVER

Delivery Pipeline
guestbook-delivery-pip...

Configured



Eclipse Orion Web IDE

Configured

guestbook-delivery-pipeline | Delivery Pipeline

INFO New! Pull Request Builds: You can now run a pipeline from a pull request or merge request. [Learn more.](#)

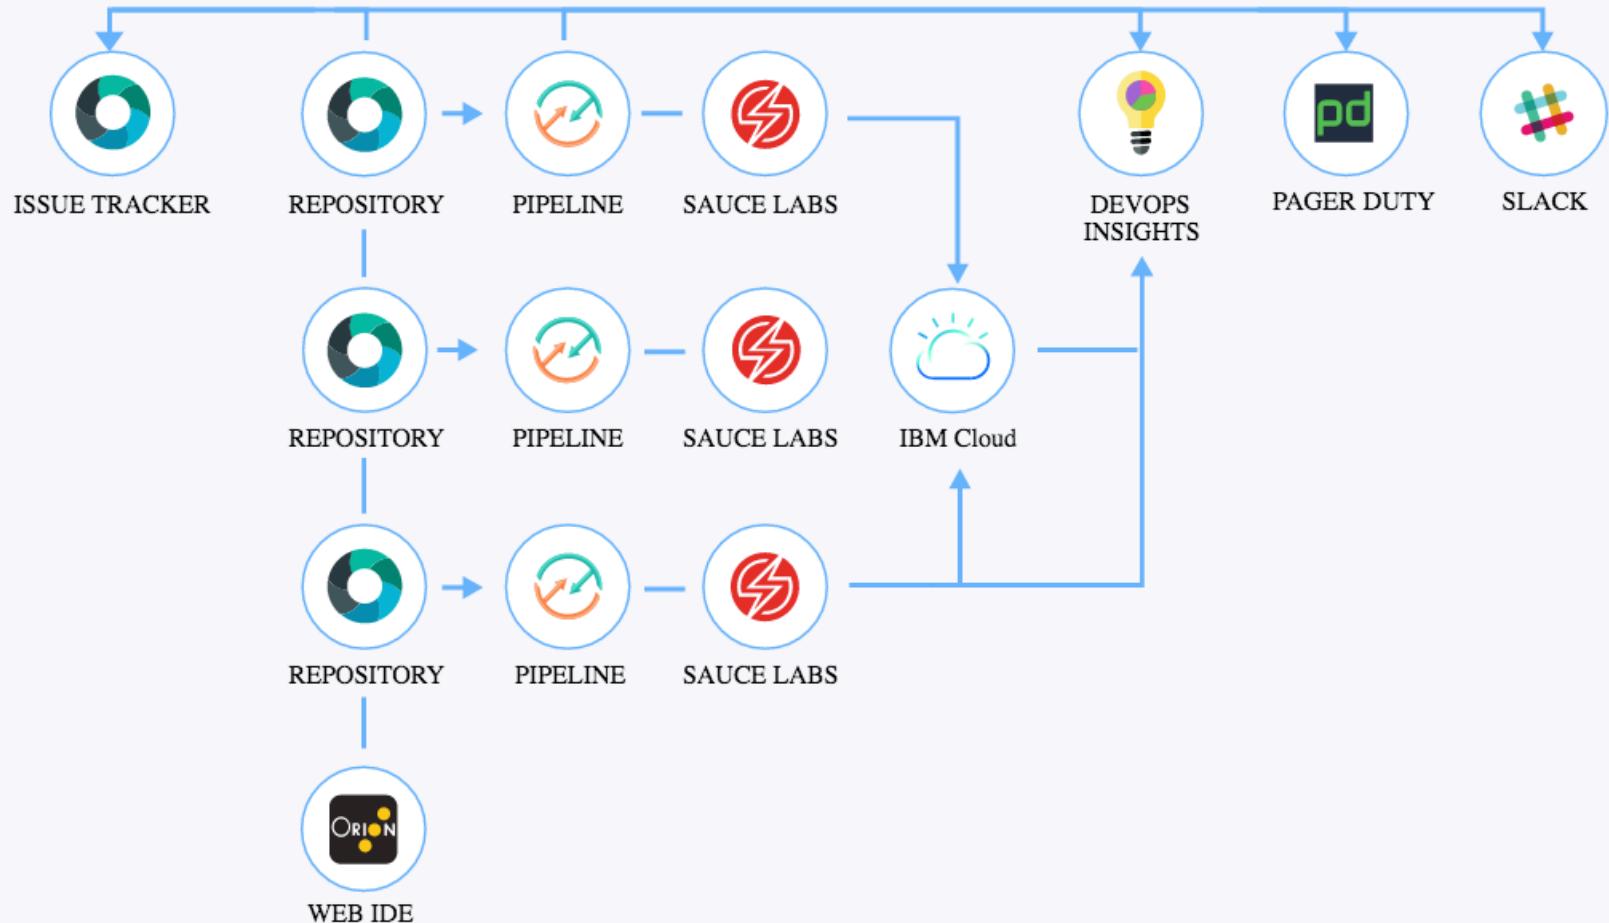
The screenshot displays a delivery pipeline interface with three stages: BUILD, CONTAINERIZE, and DEPLOY. Each stage has a green header bar indicating it is passed. The BUILD stage shows a last commit by Lin Sun from 127d ago. The CONTAINERIZE stage shows four jobs: Fetch code 1, Check dockerfile, Check registry, Build container image, and Check vulnerabilities, all passed. The DEPLOY stage shows two jobs: Deploy to Kubernetes and Check health, both passed. Each stage has a 'View logs and history' link and a 'LAST EXECUTION RESULT' section.

STAGE PASSED
LAST INPUT Last commit by Lin Sun 127d ago Merge pull request #70 from timroster/redis-slave match...
JOBSS Fetch code Passed 22d ago Unit Tests Passed 22d ago
LAST EXECUTION RESULT Fetch code 1

STAGE PASSED
LAST INPUT Stage: BUILD / Job: Fetch code Fetch code 1
JOBSS Check dockerfile Passed 22d ago Check registry Passed 22d ago Build container image Passed 22d ago Check vulnerabilities Passed 22d ago
LAST EXECUTION RESULT Check registry 3 Build container image 3

STAGE PASSED
LAST INPUT Stage: CONTAINERIZE / Job: Buil...
JOBSS Deploy to Kubernetes Passed 18d ago Check health Passed 18d ago
LAST EXECUTION RESULT No results

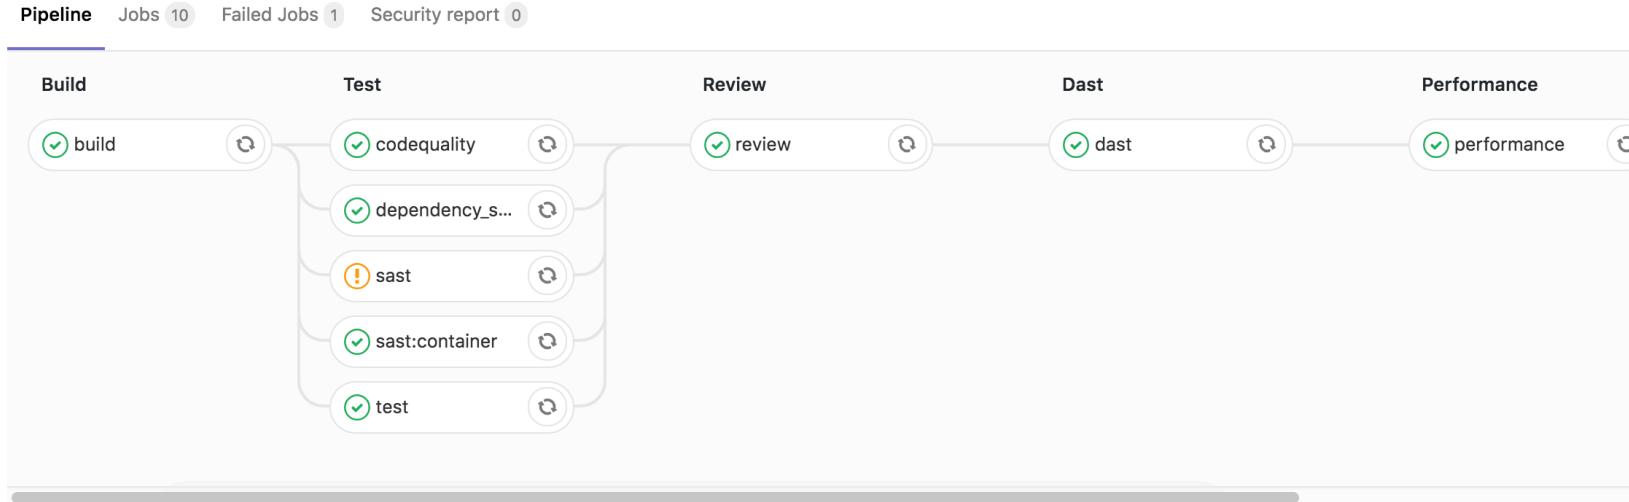
THINK CODE DELIVER RUN LEARN MANAGE CULTURE





Gitlab Auto DevOps

- Uses Heroku buildpacks,
- Build-in pipeline and configuration for Git and Kubernetes,
- Includes code quality, vulnerability, licenses, security checks, buildpack tests



Tekton

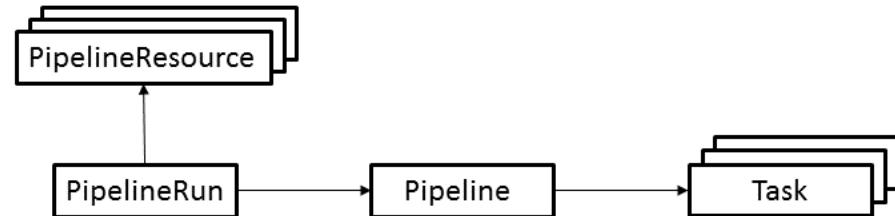
Tekton



Tekton is open-source project, born out of [Knative](#) Build, and part of the [CD Foundation](#), a [Linux Foundation](#) project. The Tekton Pipelines project has support and active commitment from companies, including IBM, Red Hat, Google, and CloudBees.

Tekton

- A **PipelineResource** defines an object that is an input (such as a git repository) or an output (such as a docker image) of the pipeline.
- A **PipelineRun** defines an execution of a pipeline. It references the **Pipeline** to run and the **PipelineResources** to use as inputs and outputs.
- A **Pipeline** defines the set of **Tasks** that compose a pipeline.
- A **Task** defines a set of build steps such as compiling code, running tests, and building and deploying images.



Tekton

```
apiVersion: tekton.dev/v1alpha1
kind: EventListener
metadata:
  name: hello-listener
spec:
  triggers:
    - binding:
        name: hello-trigger-binding
  template:
    name: hello-trigger-template
```

```
apiVersion: tekton.dev/v1alpha1
kind: TriggerTemplate
metadata:
  name: hello-trigger-template
spec:
  resources:
    - name: source-repo
      resourceRef:
        name: source-repo
        type: git
  resourceTemplates:
    - apiVersion: tekton.dev/v1alpha1
      kind: PipelineRun
      metadata:
        name: pipelinerun-$(uid)
      spec:
        pipelineRef:
          name: hello-pipeline
      resources:
        - name: source-repo
          resourceRef:
            name: source-repo
```

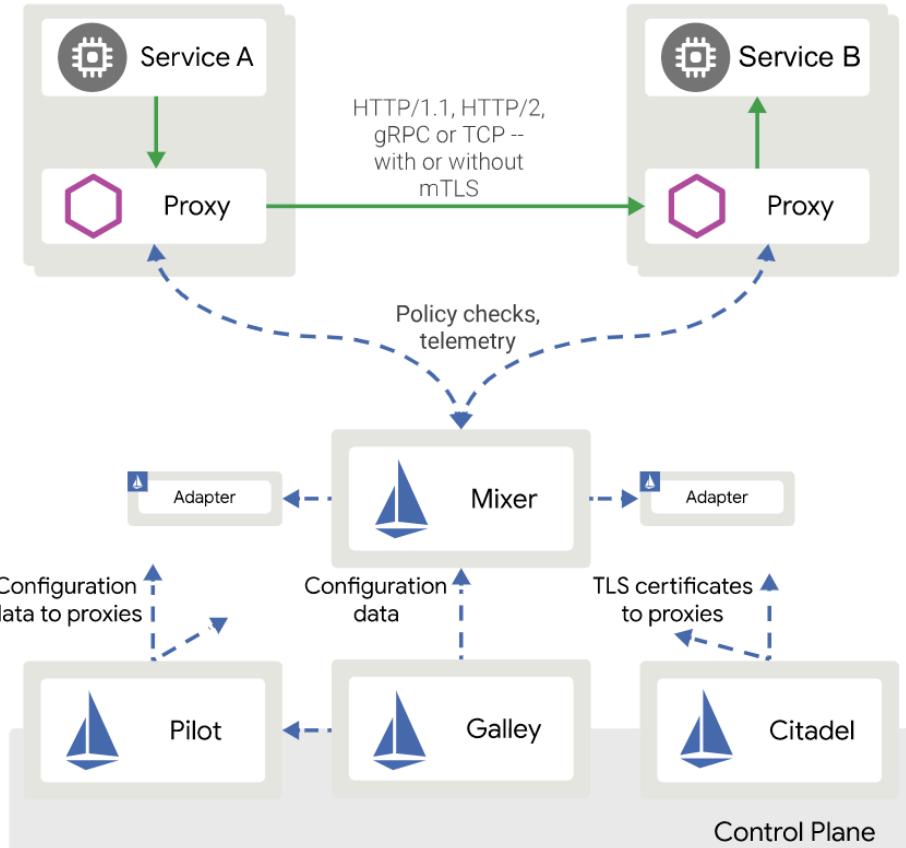
Istio

Istio simplifies configuration of :

- circuit breakers,
- timeouts,
- retries,
- A/B testing with Traffic Shifts,
- canary rollouts,
- staged rollouts with percentage-based traffic splits,
- failure recovery features.

Istio's traffic management model relies on the Envoy proxies that are deployed with your services. All traffic (data plane traffic) is proxied through Envoy, making it easy to direct and control traffic around your mesh.

The Istio API is specified using Kubernetes custom resource definitions (CRDs),



Traffic Shifts with VirtualService

```
$ kubectl apply -f samples/bookinfo/networking/virtual-service-reviews-50-v3.yaml
```

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            end-user:
              exact: jason
  route:
    - destination:
        host: reviews
        subset: v2
    - route:
        - destination:
            host: reviews
            subset: v3
```

```
      - match:
          - headers:
              user-agent:
                regex: '^.*(Android|iPhone).*$'
```

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - route:
        - destination:
            host: reviews
            subset: v1
            weight: 50
        - destination:
            host: reviews
            subset: v3
            weight: 50
```

Lab Time